



西安交通大学
XI'AN JIAOTONG UNIVERSITY



计算机系统导论

李昊

西安交通大学计算机学院

智能网络与网络安全教育部重点实验室

教学团队



李昊 副教授

hao.li@xjtu.edu.cn

兴庆校区 彭康楼234

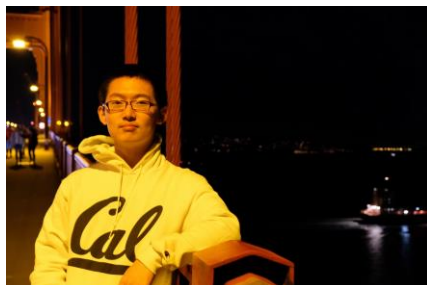


单丹枫 副教授

dfshan@xjtu.edu.cn

创新港 泓理楼6078

助教团队



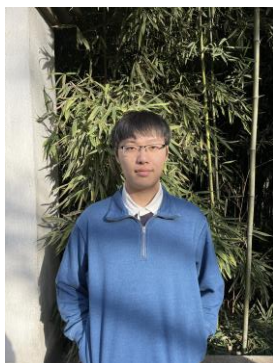
党一涵 研二 助教组长

USENIX NSDI 学生一作
研究生学分绩3/157
本科总学分绩3/32



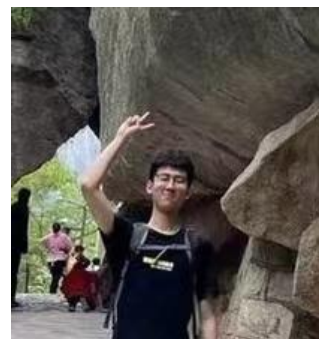
杨舜磊 研一

本科总学分绩6/171
校级一等奖学金
深交所奖学金



夏泽 研0

2次ACM ICPC银牌
1次ACM CCPC银牌
CCP认证400分



李云广 研0

本科前三年学分绩90+

教材与参考书

核心教材与原始Slides

深入理解计算机系统 第三版

Computer Systems: A Programmer's Perspective 3rd

教材与参考书

核心教材与原始Slides

深入理解计算机系统 第三版

Computer Systems: A Programmer's Perspective 3rd



教材与参考书

核心教材与原始Slides

深入理解计算机系统 第三版

Computer Systems: A Programmer's Perspective 3rd



IT图书领域的奇迹

40余国家的400余所高校将本书作为教材

哈佛大学、卡内基-梅隆大学、纽约大学、波斯顿大学、加州理工学院、加拿大国立大学、新加坡国立大学、北大、清华、复旦、上海交大、东京大学

亚洲

中国、韩国、日本、越南、老挝、柬埔寨、泰国、马来西亚、文莱、新加坡、印度尼西亚、尼泊尔、不丹、印度、巴基斯坦、斯里兰卡、伊朗、以色列、黎巴嫩、沙特阿拉伯等

欧洲

芬兰、瑞典、挪威、丹麦、俄罗斯、德国、瑞士、英国、法国、意大利、冰岛、波兰、荷兰等

大洋洲

澳大利亚、新西兰等

非洲

埃及、南非、苏丹、利比亚等

南美洲

哥伦比亚、秘鲁、巴西等

北美洲

美国、加拿大、墨西哥、哥斯达黎加等

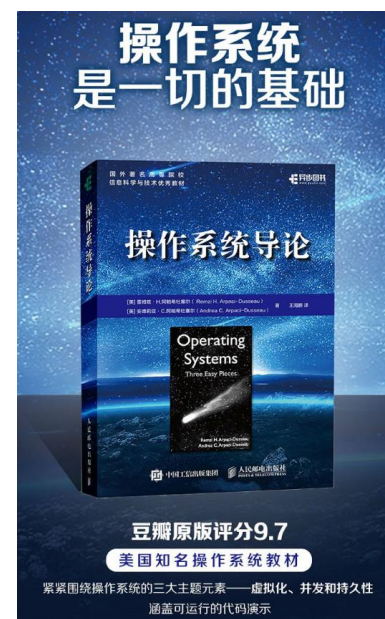
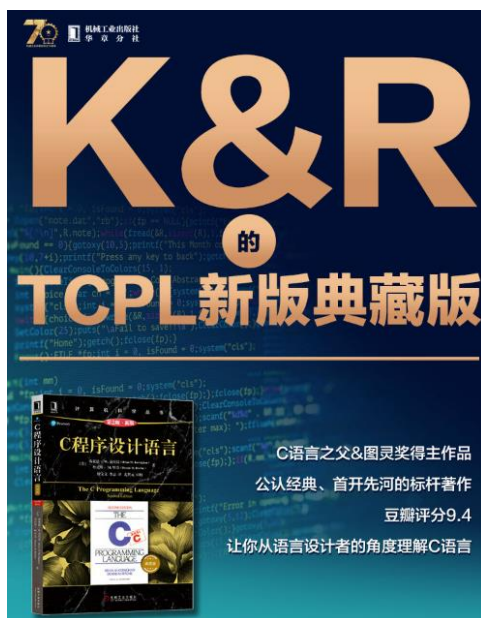


教材与参考书

参考材料

C语言程序设计 第二版 (K&R)

操作系统导论 (OSTEP)



教学计划 – 授课

内容	参考学时
信息的处理与表示	4
程序的机器级表示	8
存储器体系结构	4
程序优化	4
程序链接	4
异常控制流	4
虚拟存储器	6
系统级I/O	4
网络编程	4
并发编程	4

教学计划 – 授课

内容	参考学时
信息的处理与表示	4
程序的机器级表示	8
存储器体系结构	4
程序优化	4
程序链接	4
异常控制流	4
虚拟存储器	6
系统级I/O	4
网络编程	4
并发编程	4



教学计划 – 授课

内容	参考学时	
信息的处理与表示	4	
程序的机器级表示	8	
存储器体系结构	4	
程序优化	4	
程序链接	4	
异常控制流	4	
虚拟存储器	6	
系统级I/O	4	
网络编程	4	
并发编程	4	

教学计划 – 授课

内容	参考学时	
信息的处理与表示	4	
程序的机器级表示	8	
存储器体系结构	4	
程序优化	4	
程序链接	4	
异常控制流	4	
虚拟存储器	6	
系统级I/O	4	
网络编程	4	
并发编程	4	

教学规划 – 项目实践 (Lab)

内容	参考学时
datalab	1
bomblab	1
attacklab	1
cachelab	1
linkerlab	8
netlab	4

考核方法

考核方法

平时成绩 **10%**

考勤、课堂纪律、上课回答问题、课堂测验

考核方法

平时成绩 **10%**

考勤、课堂纪律、上课回答问题、课堂测验

项目实践 **50%**

Auto-Grading系统对代码自动打分

Anti-Cheating系统自动检测代码抄袭

抄袭是**高压线**，一经核实双方项目实践分数为0

考核方法

平时成绩 10%

考勤、课堂纪律、上课回答问题、课堂测验

项目实践 50%

Auto-Grading系统对代码自动打分

Anti-Cheating系统自动检测代码抄袭

抄袭是**高压线**，一经核实双方项目实践分数为0

期末考试 40%

以课堂和Lab内容为主

The Big Picture

Course Theme:

(Systems) Knowledge is Power!

■ Systems Knowledge

- How hardware (processors, memories, disk drives, network infrastructure) plus software (operating systems, compilers, libraries, network protocols) combine to support the execution of application programs
- How you as a programmer can best use these resources

■ Useful outcomes from taking XJTU-ICS

- Become more effective programmers
 - Able to find and eliminate bugs efficiently
 - Able to understand and tune for program performance
- Prepare for later “systems” classes in CS, ECE, INI, ...
 - Compilers, Operating Systems, Networks, Computer Architecture, Embedded Systems, Storage Systems, Computer Security, etc.

It's Important to Understand How Things Work

■ Why do I need to know this stuff?

- Abstraction is good, but don't forget reality

■ Most CS courses emphasize abstraction

- (CE courses less so)
- Abstract data types
- Asymptotic analysis

■ These abstractions have limits

- Especially in the presence of bugs
- Need to understand details of underlying implementations
- Sometimes the abstract interfaces don't provide the level of control or performance you need

Great Reality #1:

Ints are not Integers, Floats are not Reals

Great Reality #1:

Ints are not Integers, Floats are not Reals

- Example 1: Is $x^2 \geq 0$?

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

■ Float's: Yes!

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

- Float's: Yes!

- Int's:

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

- Float's: Yes!

- Int's:

- $40000 * 40000 \rightarrow 1600000000$

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

- Float's: Yes!

- Int's:

- $40000 * 40000 \rightarrow 1600000000$
- $50000 * 50000 \rightarrow ?$

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

- Float's: Yes!

- Int's:

- $40000 * 40000 \rightarrow 1600000000$

- $50000 * 50000 \rightarrow ?$

■ Example 2: Is $(x + y) + z = x + (y + z)$?

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

- Float's: Yes!

- Int's:

- $40000 * 40000 \rightarrow 1600000000$

- $50000 * 50000 \rightarrow ?$

■ Example 2: Is $(x + y) + z = x + (y + z)$?

- Unsigned & Signed Int's: Yes!

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

- Float's: Yes!

- Int's:

- $40000 * 40000 \rightarrow 1600000000$
- $50000 * 50000 \rightarrow ?$

■ Example 2: Is $(x + y) + z = x + (y + z)$?

- Unsigned & Signed Int's: Yes!
- Float's:

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

- Float's: Yes!

- Int's:

- $40000 * 40000 \rightarrow 1600000000$
- $50000 * 50000 \rightarrow ?$

■ Example 2: Is $(x + y) + z = x + (y + z)$?

- Unsigned & Signed Int's: Yes!
- Float's:
 - $(1e20 + -1e20) + 3.14 \rightarrow 3.14$

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

- Float's: Yes!

- Int's:

- $40000 * 40000 \rightarrow 1600000000$
- $50000 * 50000 \rightarrow ?$

■ Example 2: Is $(x + y) + z = x + (y + z)$?

- Unsigned & Signed Int's: Yes!

- Float's:

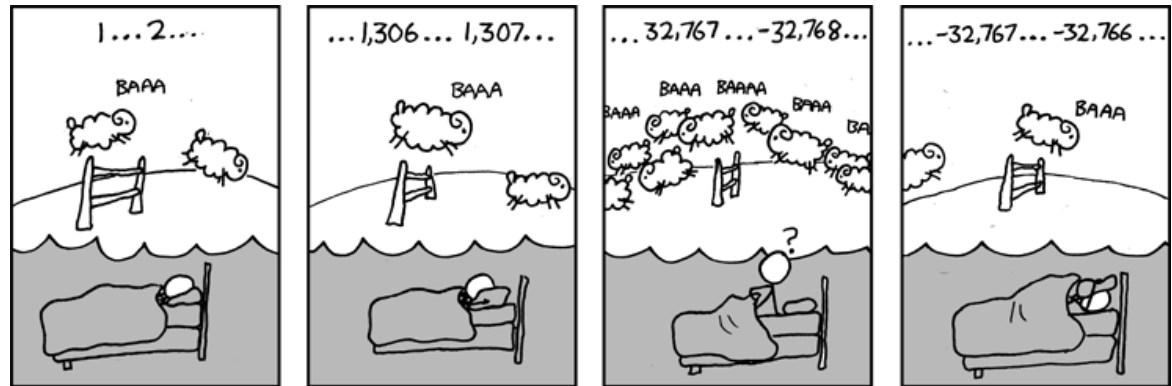
- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

■ Float's: Yes!



■ Int's:

- $40000 * 40000 \rightarrow 1600000000$
- $50000 * 50000 \rightarrow ?$

■ Example 2: Is $(x + y) + z = x + (y + z)$?

■ Unsigned & Signed Int's: Yes!

■ Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

Computer Arithmetic

■ Does not generate random values

- Arithmetic operations have important mathematical properties

■ Cannot assume all “usual” mathematical properties

- Due to finiteness of representations
- Integer operations satisfy “ring” properties
 - Commutativity, associativity, distributivity
- Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs

■ Observation

- Need to understand which abstractions apply in which contexts
- Important issues for compiler writers and serious application programmers

Great Reality #2:

You've Got to Know Assembly

- **Chances are, you'll never write programs in assembly**
 - Compilers are much better & more patient than you are
- **But: Understanding assembly is key to machine-level execution model**
 - Behavior of programs in presence of bugs
 - High-level language models break down
 - Tuning program performance
 - Understand optimizations done / not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting malware
 - x86 assembly is the language of choice!

Great Reality #3: Memory Matters

Random Access Memory Is an Unphysical Abstraction

■ Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated

■ Memory referencing bugs especially pernicious

- Effects are distant in both time and space

■ Memory performance is not uniform

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

- Result is system specific

Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    volatile struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741824; /* Possibly out of bounds */
    return s.d;
}
```

- Result is system specific

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

fun(0) --> 3.14

- Result is system specific

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

fun(0) --> 3.14

fun(1) --> 3.14

- Result is system specific

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

```
fun(0) -->      3.14  
fun(1) -->      3.14  
fun(2) -->      3.13999998664856
```

- Result is system specific

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

```
fun(0) -->      3.14  
fun(1) -->      3.14  
fun(2) -->      3.13999998664856  
fun(3) -->      2.000000061035156
```

- Result is system specific

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

```
fun(0) -->      3.14  
fun(1) -->      3.14  
fun(2) -->      3.13999998664856  
fun(3) -->      2.000000061035156  
fun(4) -->      3.14
```

- Result is system specific

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

```
fun(0) -->      3.14  
fun(1) -->      3.14  
fun(2) -->      3.13999998664856  
fun(3) -->      2.000000061035156  
fun(4) -->      3.14  
fun(6) -->      Segmentation fault
```

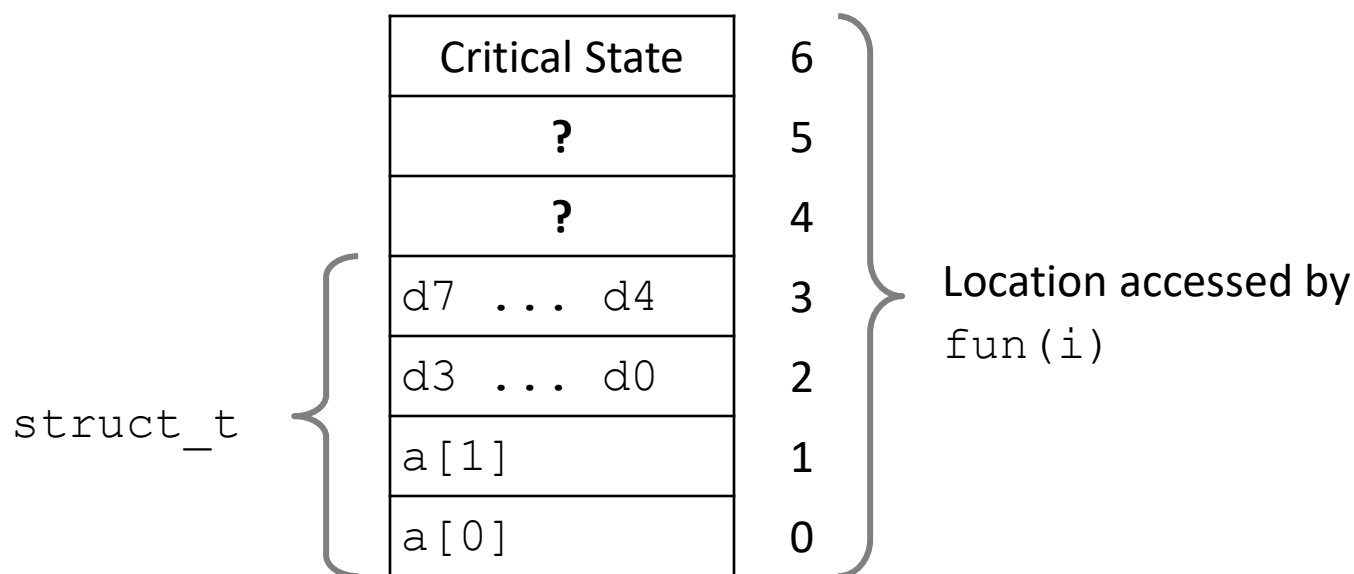
- Result is system specific

Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;
```

fun(0)	-->	3.14
fun(1)	-->	3.14
fun(2)	-->	3.1399998664856
fun(3)	-->	2.00000061035156
fun(4)	-->	3.14
fun(6)	-->	Segmentation fault

Explanation:



Memory Referencing Errors

■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

■ Can lead to nasty bugs

- Whether or not bug has any effect depends on system and compiler
- Action at a distance
 - Corrupted object logically unrelated to one being accessed
 - Effect of bug may be first observed long after it is generated

■ How can I deal with this?

- Program in Java, Ruby, Python, ML, ...
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. Valgrind)

Great Reality #4: There's more to performance than asymptotic complexity

- **Constant factors matter too!**
- **And even exact op count does not predict performance**
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- **Must understand system to optimize performance**
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Memory System Performance Example

```
void copyij(int src[2048][2048],  
            int dst[2048][2048])  
{  
    int i,j;  
    for (i = 0; i < 2048; i++)  
        for (j = 0; j < 2048; j++)  
            dst[i][j] = src[i][j];  
}
```

Memory System Performance Example


```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```



Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

4.3ms

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

81.8ms

2.0 GHz Intel Core i7 Haswell

Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

4.3ms

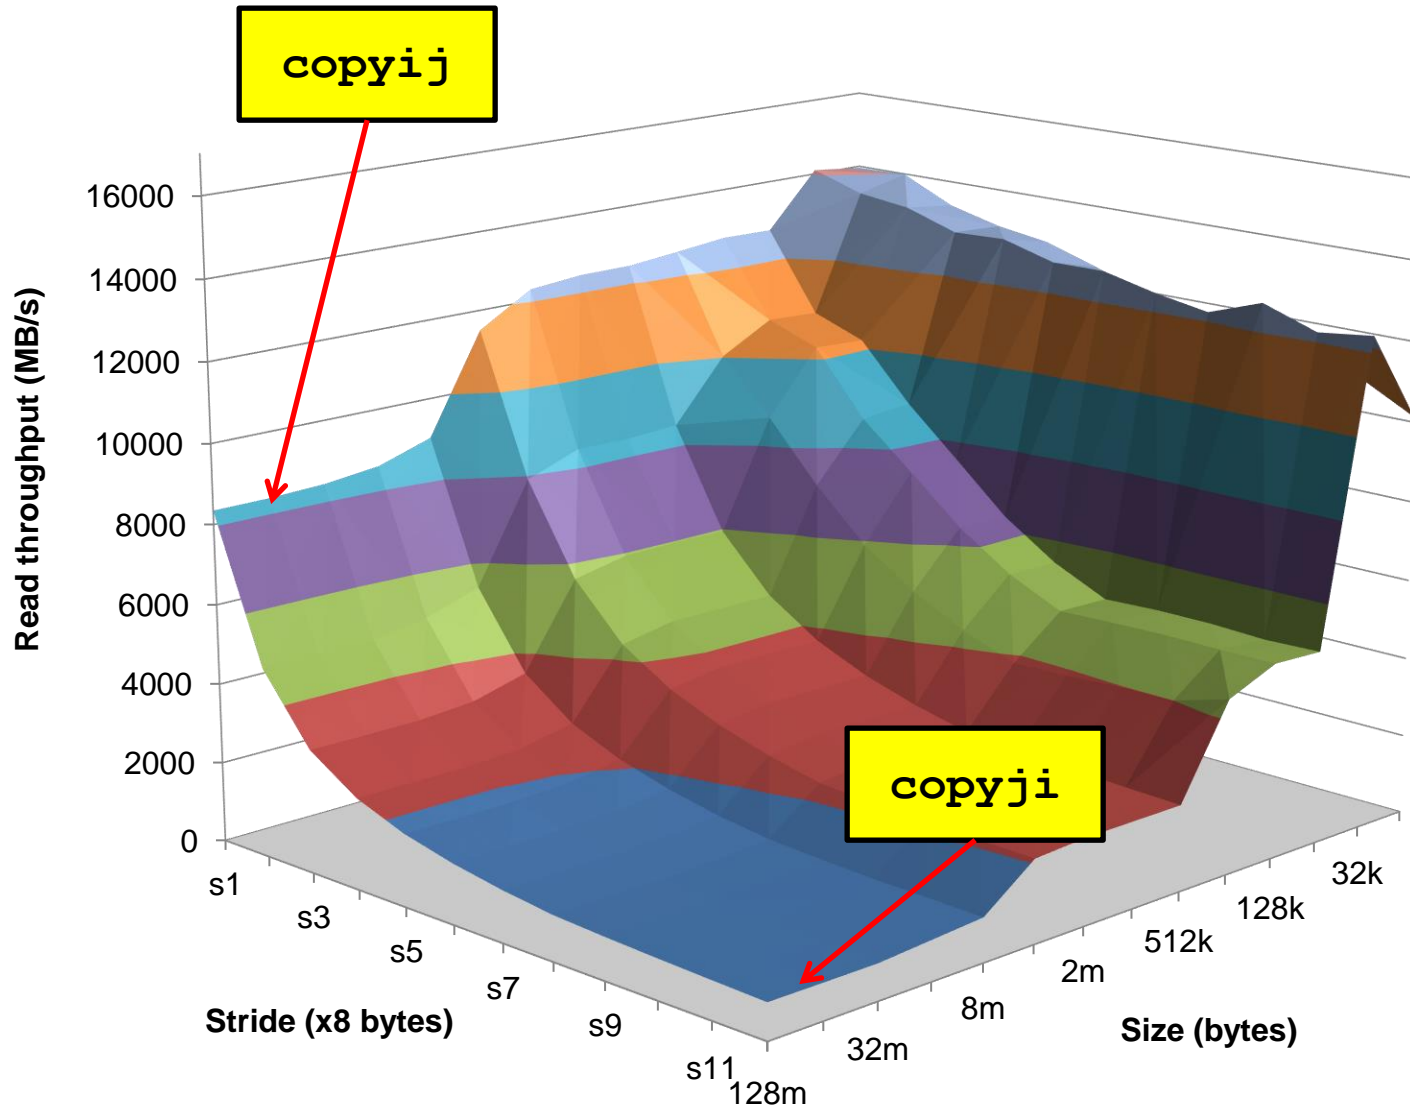
```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

81.8ms

2.0 GHz Intel Core i7 Haswell

- Hierarchical memory organization
- Performance depends on access patterns
 - Including how step through multi-dimensional array

Why The Performance Differs



Great Reality #5:

Computers do more than execute programs

- **They need to get data in and out**
 - I/O system critical to program reliability and performance

- **They communicate with each other over networks**
 - Many system-level issues arise in presence of network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues

Course Perspective

■ Most Systems Courses are Builder-Centric

- Computer Architecture
 - Design pipelined processor in Verilog
- Operating Systems
 - Implement sample portions of operating system
- Compilers
 - Write compiler for simple language
- Networking
 - Implement and simulate network protocols

Course Perspective (Cont.)

■ Our Course is Programmer-Centric

- By knowing more about the underlying system, you can be more effective as a programmer
- Enable you to
 - Write programs that are more reliable and efficient
 - Incorporate features that require hooks into OS
 - E.g., concurrency, signal handlers
- Cover material in this course that you won't see elsewhere
- Not just a course for dedicated hackers
 - **We bring out the hidden hacker in everyone!**

*Welcome
and Enjoy!*